


What is an operating system?


- Layer between applications and hardware



- Makes hardware useful to the programmer
- [Usually] Provides abstractions for applications
 - Manages and hides details of hardware
 - Accesses hardware through low/level interfaces unavailable to applications
- [Often] Provides protection
 - Prevents one process/user from clobbering another


Primitive Operating Systems

- Just a library of standard services [no protection]




- Standard interface above hardware-specific drivers, etc.
- **Simplifying assumptions**
 - System runs one program at a time
 - No bad users or programs (often bad assumption)
- **Problem: Poor utilization**
 - ...of hardware (e.g., CPU idle while waiting for disk)
 - ...of human user (must wait for each program to finish)

Multitasking




- **Idea: Run more than one process at once**
 - When one process blocks (waiting for disk, network, user input, etc.) run another process
- **Problem: What can ill-behaved process do?**

Multitasking




- **Idea: Run more than one process at once**
 - When one process blocks (waiting for disk, network, user input, etc.) run another process
- **Problem: What can ill-behaved process do?**
 - Go into infinite loop and never relinquish CPU
 - Scribble over other processes' memory to make them fail
- **OS provides mechanisms to address these problems**
 - *Preemption* – take CPU away from looping process
 - *Memory protection* – protect process's memory from one another

Multi-user OSes



- Many OSes use *protection* to serve distrustful users
- **Idea: With N users, system not N times slower**
 - Users' demands for CPU, memory, etc. are bursty
 - Win by giving resources to users who actually need them
- **What can go wrong?**

Multi-user OSes

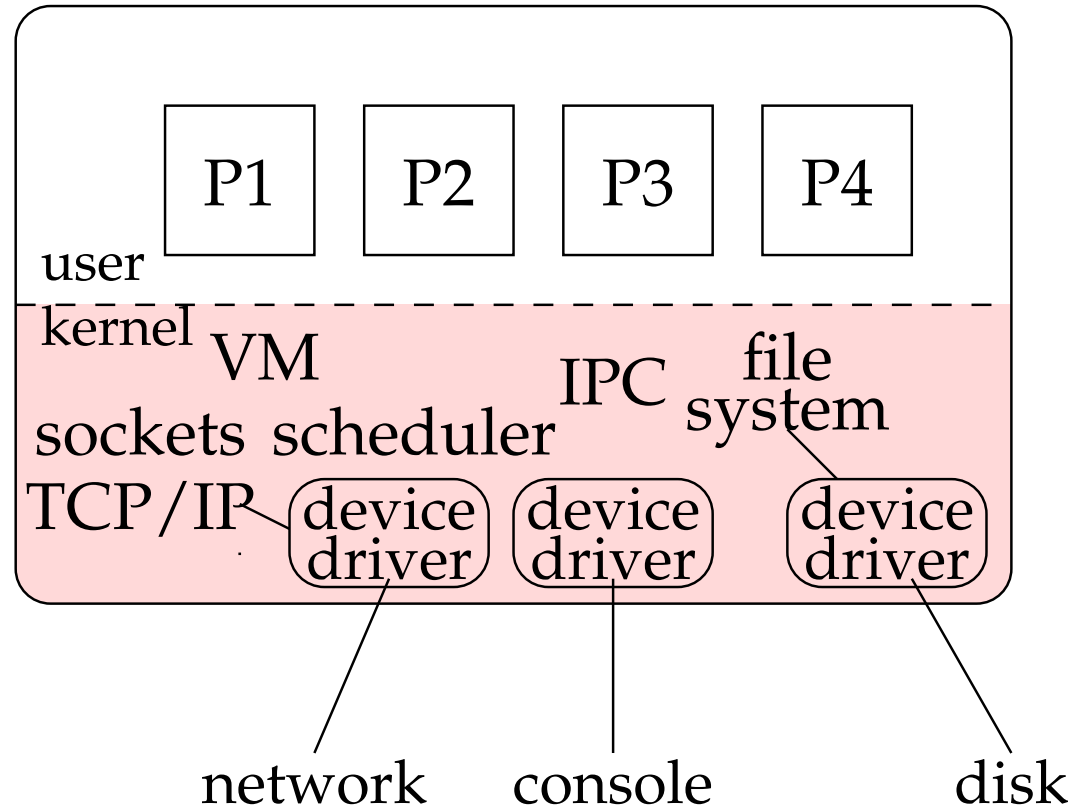


- Many OSes use *protection* to serve distrustful users
- **Idea: With N users, system not N times slower**
 - Users' demands for CPU, memory, etc. are bursty
 - Win by giving resources to users who actually need them
- **What can go wrong?**
 - Users are gluttons, use too much CPU, etc. (need policies)
 - Total memory usage greater than in machine (must virtualize)
 - Super-linear slowdown with increasing demand (thrashing)

Protection


- **Mechanisms that isolate bad programs and people**
- **Pre-emption:**
 - Give application a resource, take it away if needed elsewhere
- **Interposition/mediation:**
 - Place OS between application and “stuff”
 - Track all pieces that application allowed to use (e.g., in table)
 - On every access, look in table to check that access legal
- **Privileged & unprivileged modes in CPUs :**
 - Applications unprivileged (user/unprivileged mode)
 - OS privileged (privileged/supervisor mode)
 - Protection operations can only be done in privileged mode

Typical OS structure



- Most software runs as user-level processes (P[1-4])
- OS *kernel* runs in *privileged* mode **[shaded]**
 - Creates/deletes processes
 - Provides access to hardware

System calls

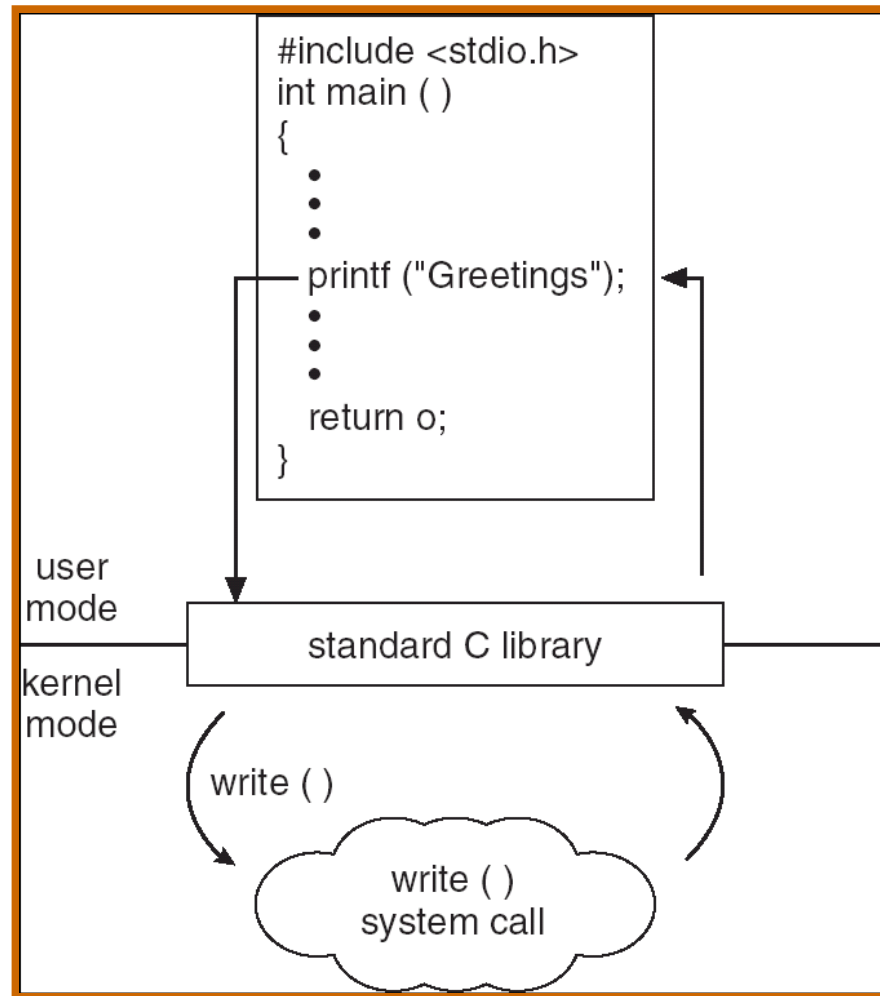


- Applications can invoke kernel through *system calls*
 - Special instruction transfers control to kernel
 - ...which dispatches to one of few hundred syscall handlers

System calls (continued)

- **Goal: Do things app. can't do in unprivileged mode**
 - Like a library call, but into more privileged kernel code
- **Kernel supplies well-defined *system call* interface**
 - Applications set up syscall arguments and *trap* to kernel
 - Kernel performs operation and returns result
- **Higher-level functions built on syscall interface**
 - `printf`, `scanf`, `gets`, etc. all user-level code
- **Example: POSIX/UNIX interface**
 - `open`, `close`, `read`, `write`, ...

System call example



- **Standard library implemented in terms of syscalls**
 - *printf* – in libc, has same privileges as application
 - calls *write* – in kernel, which can send bits out serial port

UNIX file system calls

- **Applications “open” files (or devices) by name**
 - I/O happens through open files
- `int open(char *path, int flags, /*mode*/...);`
 - flags: `O_RDONLY`, `O_WRONLY`, `O_RDWR`
 - `O_CREAT`: create the file if non-existent
 - `O_EXCL`: (w. `O_CREAT`) create if file exists already
 - `O_TRUNC`: Truncate the file
 - `O_APPEND`: Start writing from end of file
 - mode: final argument with `O_CREAT`
- **Returns file descriptor—used for all I/O to file**

Error returns

- **What if open fails? Returns -1 (invalid fd)**
- **Most system calls return -1 on failure**
 - Specific kind of error in global int errno
- **#include <sys/errno.h> for possible values**
 - 2 = ENOENT “No such file or directory”
 - 13 = EACCES “Permission Denied”
- **perror function prints human-readable message**
 - `perror ("initfile");`
→ “initfile: No such file or directory”

Operations on file descriptors

- `int read (int fd, void *buf, int nbytes);`
 - Returns number of bytes read
 - Returns 0 bytes at end of file, or -1 on error
- `int write (int fd, void *buf, int nbytes);`
 - Returns number of bytes written, -1 on error
- `off_t lseek (int fd, off_t pos, int whence);`
 - whence: 0 – start, 1 – current, 2 – end
 - ▷ Returns previous file offset, or -1 on error
- `int close (int fd);`

File descriptor numbers

- **File descriptors are inherited by processes**
 - When one process spawns another, same fds by default
- **Descriptors 0, 1, and 2 have special meaning**
 - 0 – “standard input” (stdin in ANSI C)
 - 1 – “standard output” (stdout, printf in ANSI C)
 - 2 – “standard error” (stderr, perror in ANSI C)
 - Normally all three attached to terminal
- **Example: type.c**
 - Prints the contents of a file to stdout

type.c

```
void
typefile (char *filename)
{
    int fd, nread;
    char buf[1024];

    fd = open (filename, O_RDONLY);
    if (fd == -1) {
        perror (filename);
        return;
    }

    while ((nread = read (fd, buf, sizeof (buf))) > 0)
        write (1, buf, nread);

    close (fd);
}
```


Different system contexts

- A system can typically be in one of several contexts
- *User-level* – running an application
- **Kernel process context**
 - Running kernel code on behalf of a particular process
 - E.g., performing system call
 - Also exception (mem. fault, numeric exception, etc.)
 - Or executing a kernel-only process (e.g., network file server)
- **Kernel code not associated w. a process**
 - Timer interrupt (hardclock)
 - Device interrupt
 - “Softirqs”, “Tasklets” (Linux-specific terms)
- **Context switch code – changing address spaces**

CPU preemption

- **Protection mechanism to prevent monopolizing CPU**
- **E.g., kernel programs timer to interrupt every 10 ms**
 - Must be in supervisor mode to write appropriate I/O registers
 - User code cannot re-program interval timer
- **Kernel sets interrupt to vector back to kernel**
 - Regains control whenever interval timer fires
 - Gives CPU to another process if someone else needs it
 - Note: must be in supervisor mode to set interrupt entry points
 - No way for user code to hijack interrupt handler
- **Result: Cannot monopolize CPU with infinite loop**
 - At worst get $1/N$ of CPU with N CPU-hungry processes

Protection is not security

- How *can* you monopolize CPU?

Protection is not security

- How *can* you monopolize CPU?
- Use multiple processes
- For many years, could wedge most OSes with

```
int main() { while(1) fork(); }
```

 - Keeps creating more processes until system out of proc. slots
- Other techniques: use all memory (chill program)
- Typically solved with technical/social combination
 - Technical solution: Limit processes per user
 - Social: Reboot and yell at annoying users
 - Social: Pass laws (often debatable whether a good idea)

Address translation

- **Protect mem. of one program from actions of another**
- **Definitions**
 - *Address space*: all memory locations a program can name
 - *Virtual address*: addresses in process' address space
 - *Physical address*: address of real memory
 - *Translation*: map virtual to physical addresses
- **Translation done on every load and store**
 - Modern CPUs do this in hardware for speed
- **Idea: If you can't name it, you can't touch it**
 - Ensure one process's translations don't include any other process's memory

Resource allocation & performance

- **Multitasking permits higher resource utilization**
- **Simple example:**
 - Process downloading large file mostly waits for network
 - You play a game while downloading the file
 - Higher CPU utilization than if just downloading
- **Complexity arises with cost of switching**
- **Example: Say disk 1,000 times slower than memory**
 - 1 GB memory in machine
 - 2 Processes want to run, each use 1 GB
 - Can switch processes by swapping them out to disk
 - Faster to run one at a time than keep context switching

Useful properties to exploit

- **Skew**

- 80% of time taken by 20% of code
- 10% of memory absorbs 90% of references
- Basis behind cache: place 10% in fast memory, 90% in slow, usually looks like one big fast memory

- **Past predicts future (a.k.a. temporal locality)**

- What's the best cache entry to replace?
- If past = future, then least-recently-used entry

- **Note conflict between fairness & throughput**

- Higher throughput (fewer cache misses, etc.) to keep running same process
- But fairness says should periodically preempt CPU and give it to next process